

# Fun With Graphite

# What is Graphite?

- Time-series based metrics DB
- Graphing engine for stored metrics

# Time-Series Metrics

- Metric data points are expected at a given frequency
- Insufficient metrics in an interval mean nothing is stored
  - This can be adjusted

# How Are Metrics Stored?

- Whisper files
- Each metric is a separate file
- Whisper files are pre-allocated
- Files can contain multiple aggregations

# Whisper File Location

- /opt/graphite/storage/whisper
- Convert '.' to '/' in path
  - servers.hostname.loadavg.01
  - /opt/graphite/storage/whisper/  
servers/hostname/loadavg/01.wsp

# Whisper File Contents

- Pre-allocated
  - Size doesn't change
  - All timestamps pre-written
  - Only stores a timestamp & a metric result



```
$ whisper-dump.py 01.wsp
```

```
Meta data:
```

```
  aggregation method: average
```

```
  max retention: 31536000
```

```
  xFilesFactor: 0.5
```

```
Archive 0 info:
```

```
  offset: 40
```

```
  seconds per point: 30
```

```
  points: 20160
```

```
  retention: 604800
```

```
  size: 241920
```

```
Archive 1 info:
```

```
  offset: 241960
```

```
  seconds per point: 300
```

```
  points: 105120
```

```
  retention: 31536000
```

```
  size: 1261440
```

```
Archive 0 data:
```

```
0: 1379553930, 2.7700000000000000177635683940025046
```

```
1: 1379553960, 2.3199999999999998401278844539774582
```

```
2: 1379553990, 2.0400000000000000355271367880050093
```

```
...
```

```
$ whisper-dump.py 01.wsp
Meta data:
  aggregation method: average
  max retention: 31536000
  xFilesFactor: 0.5

Archive 0 info:
  offset: 40
  seconds per point: 30
  points: 20160
  retention: 604800
  size: 241920

Archive 1 info:
  offset: 241960
  seconds per point: 300
  points: 105120
  retention: 31536000
  size: 1261440

Archive 0 data:
0: 1379553930, 2.7700000000000000177635683940025046
1: 1379553960, 2.3199999999999998401278844539774582
2: 1379553990, 2.0400000000000000355271367880050093
...
```

#### Aggregation method:

- \* Average (default) – all metrics received in an interval are averaged into a single metric value for the interval
- \* Maximum, minimum – only the highest or lowest value is used – all others are discarded
- \* Sum – All metrics received during the interval are summed up and the total is stored for the interval

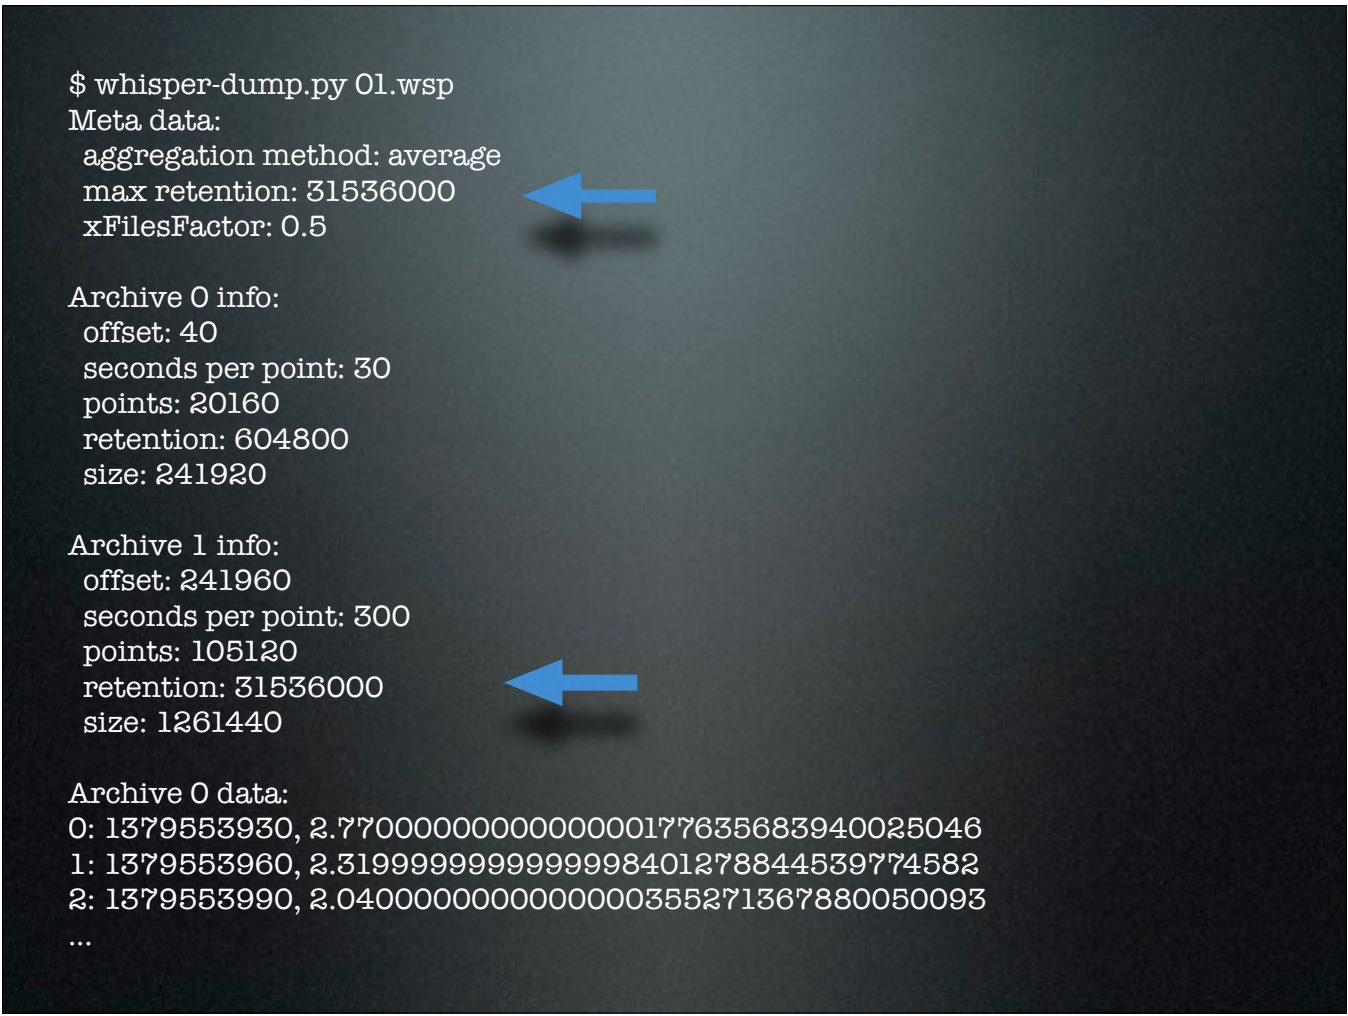


```
$ whisper-dump.py 01.wsp
Meta data:
  aggregation method: average
  max retention: 31536000
  xFilesFactor: 0.5

Archive 0 info:
  offset: 40
  seconds per point: 30
  points: 20160
  retention: 604800
  size: 241920

Archive 1 info:
  offset: 241960
  seconds per point: 300
  points: 105120
  retention: 31536000
  size: 1261440

Archive 0 data:
0: 1379553930, 2.7700000000000000177635683940025046
1: 1379553960, 2.3199999999999998401278844539774582
2: 1379553990, 2.0400000000000000355271367880050093
...
```



Max Retention:


\* Matches the longest retention of all of the archives included in this metric

```
$ whisper-dump.py 01.wsp
Meta data:
  aggregation method: average
  max retention: 31536000
  xFilesFactor: 0.5

Archive 0 info:
  offset: 40
  seconds per point: 30
  points: 20160
  retention: 604800
  size: 241920

Archive 1 info:
  offset: 241960
  seconds per point: 300
  points: 105120
  retention: 31536000
  size: 1261440

Archive 0 data:
0: 1379553930, 2.7700000000000000177635683940025046
1: 1379553960, 2.3199999999999998401278844539774582
2: 1379553990, 2.0400000000000000355271367880050093
...
```



#### xFilesFactor:


- \* Percentage of metrics per interval that must be non-null in order for the interval to be considered valid.
- \* If this is set to 0.5, the highest resolution retention is 30 seconds & metrics are injected every 3 seconds, at least 8 metrics per 30 seconds must be non-null values in order for that 30 second interval to have a value attached to it. If less than 8 metrics are received in that 30 second interval, the interval is given a null value (NaN – Not a Number).

```
$ whisper-dump.py 01.wsp
Meta data:
  aggregation method: average
  max retention: 31536000
  xFilesFactor: 0.5

Archive 0 info:
  offset: 40
  seconds per point: 30
  points: 20160
  retention: 604800
  size: 241920

Archive 1 info:
  offset: 241960
  seconds per point: 300
  points: 105120
  retention: 31536000
  size: 1261440

Archive 0 data:
0: 1379553930, 2.7700000000000000177635683940025046
1: 1379553960, 2.3199999999999998401278844539774582
2: 1379553990, 2.0400000000000000355271367880050093
...
```



#### Archives:

- \* Archives are synonymous with retention schemas or retentions
- \* A metric can have multiple retentions
- \* This particular metric has 2 retentions defined:
  - Every 30 seconds for 7 days
  - Every 5 minutes for 1 year

```
$ whisper-dump.py 01.wsp
Meta data:
  aggregation method: average
  max retention: 31536000
  xFilesFactor: 0.5

Archive 0 info:
  offset: 40
  seconds per point: 30
  points: 20160
  retention: 604800
  size: 241920

Archive 1 info:
  offset: 241960
  seconds per point: 300
  points: 105120
  retention: 31536000
  size: 1261440

Archive 0 data:
0: 1379553930, 2.7700000000000000177635683940025046
1: 1379553960, 2.3199999999999998401278844539774582
2: 1379553990, 2.0400000000000000355271367880050093
...
```

Offset:

\* Why byte in the Whisper file the archive (retention) data begins to be stored at




```
$ whisper-dump.py 01.wsp
Meta data:
  aggregation method: average
  max retention: 31536000
  xFilesFactor: 0.5

Archive 0 info:
  offset: 40
  seconds per point: 30
  points: 20160
  retention: 604800
  size: 241920

Archive 1 info:
  offset: 241960
  seconds per point: 300
  points: 105120
  retention: 31536000
  size: 1261440

Archive 0 data:
0: 1379553930, 2.7700000000000000177635683940025046
1: 1379553960, 2.3199999999999998401278844539774582
2: 1379553990, 2.0400000000000000355271367880050093
...
```



Seconds per point:

- \* How long of an interval each metric point stored represents

- \* In this file each interval takes the average of all the metrics injected during the interval and stores the result as a single value for the interval defined in each archive (or retention)




```
$ whisper-dump.py 01.wsp
Meta data:
  aggregation method: average
  max retention: 31536000
  xFilesFactor: 0.5

Archive 0 info:
  offset: 40
  seconds per point: 30
  points: 20160
  retention: 604800
  size: 241920

Archive 1 info:
  offset: 241960
  seconds per point: 300
  points: 105120
  retention: 31536000
  size: 1261440

Archive 0 data:
0: 1379553930, 2.7700000000000000177635683940025046
1: 1379553960, 2.3199999999999998401278844539774582
2: 1379553990, 2.0400000000000000355271367880050093
...
```



#### Points:

- \* The total number of points comprising the retention period
- \* 30 seconds \* 604,800 seconds (7 days) = 20,160 points
- \* 300 seconds (5 minutes) \* 31,536,000 (365 days) = 105,120 points

```
$ whisper-dump.py 01.wsp
Meta data:
  aggregation method: average
  max retention: 31536000
  xFilesFactor: 0.5

Archive 0 info:
  offset: 40
  seconds per point: 30
  points: 20160
  retention: 604800
  size: 241920
  ←

Archive 1 info:
  offset: 241960
  seconds per point: 300
  points: 105120
  retention: 31536000
  size: 1261440
  ←

Archive 0 data:
0: 1379553930, 2.7700000000000000177635683940025046
1: 1379553960, 2.3199999999999998401278844539774582
2: 1379553990, 2.0400000000000000355271367880050093
...
```

#### Retention:

\* Length of time to keep storing metrics for (in seconds)

- 604,800 seconds = 7 days
- 31,536,000 seconds = 1 year

```
$ whisper-dump.py 01.wsp
Meta data:
  aggregation method: average
  max retention: 31536000
  xFilesFactor: 0.5

Archive 0 info:
  offset: 40
  seconds per point: 30
  points: 20160
  retention: 604800
  size: 241920
  ←

Archive 1 info:
  offset: 241960
  seconds per point: 300
  points: 105120
  retention: 31536000
  size: 1261440
  ←

Archive 0 data:
0: 1379553930, 2.7700000000000000177635683940025046
1: 1379553960, 2.3199999999999998401278844539774582
2: 1379553990, 2.0400000000000000355271367880050093
...
```

#### Size:


- \* Total size of the archive (retention) in bytes
- \* Each metric stored = ~12 bytes

```
$ whisper-dump.py 01.wsp
Meta data:
  aggregation method: average
  max retention: 31536000
  xFilesFactor: 0.5

Archive 0 info:
  offset: 40
  seconds per point: 30
  points: 20160
  retention: 604800
  size: 241920

Archive 1 info:
  offset: 241960
  seconds per point: 300
  points: 105120
  retention: 31536000
  size: 1261440

Archive 0 data:
0: 1379553930, 2.7700000000000000177635683940025046
1: 1379553960, 2.3199999999999998401278844539774582
2: 1379553990, 2.0400000000000000355271367880050093
...
```



Archive data:

\* 3 fields:

- ID/counter
- Timestamp (epoch)
- Metric value

# Graphite Daemons

- carbon-relay
- carbon-cache
- carbon-aggregator



# carbon-relay

carbon-relay:

- \* Like a router for metrics
- \* Directs metrics to one or more carbon-cache daemons on one or more servers
- \* Can direct metrics based on rules or use consistent hashing with configurable number of replicas
  - Rule-based reads a configuration file with regexes defining which metrics should be sent to which carbon-caches
  - Consistent hashing will write each metric to a pool of carbon-caches replicating each metric N number of times based on what the replication is configured to be

# carbon-cache

carbon-cache:

- \* Receives the metrics and handles writing them to disk
- \* Uses memory to cache recent metrics and asynchronously (by default) writes them to disk to balance IO
- \* Follows schema configurations for retentions in storage-schema.conf to define the structure of Whisper files & determine how to store the metric results.
- \* Responsible for returning the desired metrics when requested by the API
- \* Multiple carbon-caches can run on a given server to improve IO utilization across multiple processes

# carbon-aggregator

carbon-aggregator:

- \* Not currently used at AWeber
- \* Pre-processes aggregations before sending them to carbon-cache
- \* Because results are buffered & pre-processed they aren't available in "real-time" like they are if you're using carbon-cache directly
- \* Can help improve IO load

# Configuration

- carbon.conf
- storage-schema.conf
- storage-aggregation.conf

# carbon.conf

- \* Configuration file for carbon-relay & carbon-cache daemons
- \* Defines ports & IP addresses each daemon binds to (listens on)
- \* Defines how carbon-cache creates & writes to Whisper files
- \* Defines how carbon-cache uses memory to cache metrics & how aggressively to write to disk
- \* Defines how carbon-relay is configured (rules-based or consistent hashing)
- \* Defines which carbon-cache daemons/hosts carbon-relay should send metrics to
- \* Defines how many metrics carbon-relay can queue waiting for carbon-cache to accept before either dropping new metrics injected or rejecting metrics
  - Because metric injection is via TCP by default, setting carbon-relay to reject metrics vs. dropping them can provide the ability for the client to retransmit metrics (pros/cons to each)



# storage-schema.conf

```
[servers_loadavg]  
priority = 50  
pattern = ^servers\.*\.loadavg\.*  
retentions = 30s:7d,5m:1y
```

- \* Defines the retentions (archives) each metric file will include
- \* Uses a regex to determine which metric matches the rule defined
- \* Priorities can be used to define higher granularity of specific metrics within a broader scope of metrics of similar names
- \* Retentions are defined as interval:retention period
  - Interval is how often metrics are stored
  - Retention period is the length of time a particular metric is stored for

# storage-aggregation.conf

```
[chef_handler]  
pattern = ^servers\..*chef\..*\fail$  
x_files_factor = 0  
aggregation_method = sum
```

- \* Defines how metrics received over a particular interval are aggregated & stored for that interval
  - In this example, the metric has an interval of 1 minute, so multiple injections of this metric within that 1 minute interval are summed up before being stored
- \* Sets the xFilesFactor setting for the Whisper file (defining how many results per interval are required for a valid result)
  - Because this is set to zero, *any* result is sufficient to be considered a valid metric for the interval. Only when *no* results are received for the interval (1 minute) will the interval be given a null value.
  - Because chef-client can be run at any time arbitrarily, a high resolution is used for this metric so that any chef-client run is recorded. The downside to this is that there is a lot of gaps in the file where 'null' values will be written wasting space.
    - a. Supposedly the next storage mechanism (Ceres) will allow for only storing actual metrics in expanding/contracting DB files rather than using fixed-size DB files



# API

carbon-aggregator:


- \* Not currently used at AWeber
- \* Pre-processes aggregations before sending them to carbon-cache
- \* Because results are buffered & pre-processed they aren't available in "real-time" like they are if you're using carbon-cache directly
- \* Can help improve IO load

# API Features

- URL-based
- Globbing
- Functions (nested)

# URL-based API


- Query-strings used to access metrics
- <http://graphite.colo.lair/render/?target=servers.graphite3.loadavg.01>





# Globber/Regexes for Multiple Metrics

- Globbing for multiple metrics
- [http://graphite.colo.lair/render/?target=servers.\\*.loadavg.01](http://graphite.colo.lair/render/?target=servers.*.loadavg.01)

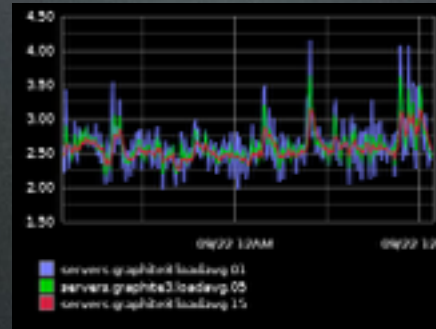


Graph represents 1-minute load average across all hosts

\* Note that the legend automatically disappears when there's >10 metrics plotted at the same time

# Globber/Regexes for Multiple Metrics


- Simple regex for multiple metrics
- <http://graphite.colo.lair/render/?target=servers.graphite3.loadavg.{01,05,15}>



This graph shows 1-minute, 5-minute & 15-minute load average for the graphite3 host.


# Using Functions

- Simple function for setting the legend names
- [http://graphite.colo.lair/render/?target=aliasByNode\(servers.graphite3.loadavg.{01,05,15},3\)](http://graphite.colo.lair/render/?target=aliasByNode(servers.graphite3.loadavg.{01,05,15},3))



# Using Functions

- Nesting multiple functions
- [http://graphite.colo.lair/render/?target=alias\(averageSeries\(servers.graphite3.loadavg.{01,05,15}\),"Average+of+all+load+average+results"\)](http://graphite.colo.lair/render/?target=alias(averageSeries(servers.graphite3.loadavg.{01,05,15}),)



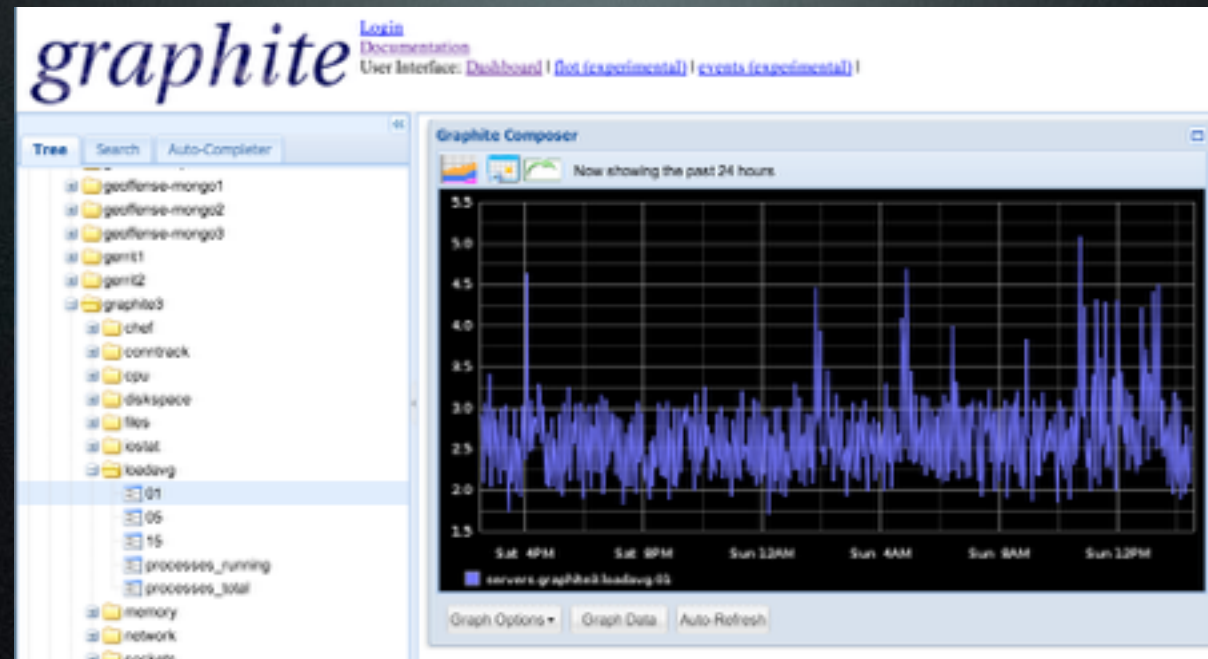
An average of all 3 metric results for loadavg (1, 5 & 15 minute) is plotted & the legend alias is modified to be more friendly.

# Web UI

- Graph Generator
- Dashboard Generator



# Graph Generator



# Graph Generator

- Pros:
  - Useful for quickly building complex graphs
  - Apply & remove functions on metrics quickly
- Cons:
  - Only allows for creating a single graph
  - Can't really save created graphs\*

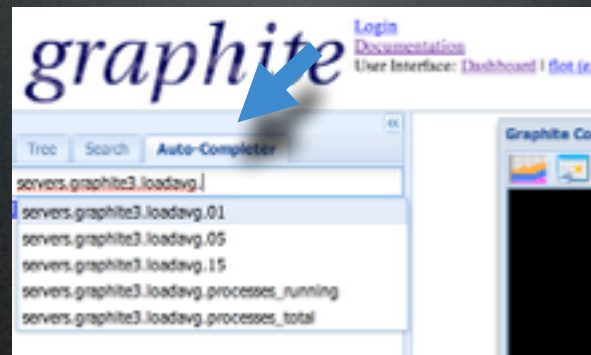
\* It is possible to save graphs, but currently only via the 'root' user (not the system's root user).

– Not a feature we use at AWeber – use the Dashboards instead

– There is LDAP integration, but other dashboards are probably preferred over spending time on this

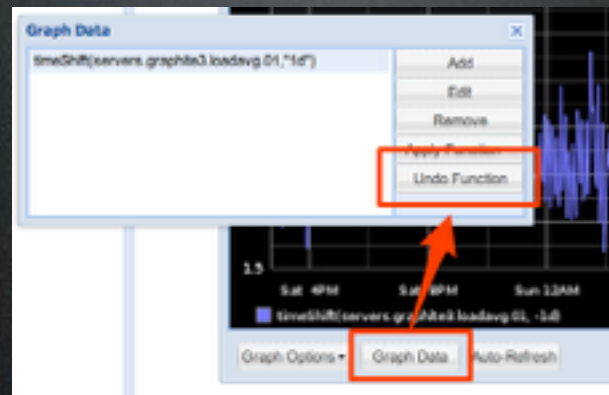
# Graph Generator Tip

- Auto-Completer tab allows for “finding” metrics
  - Start typing metric to get list
  - Must start from the beginning of the metric name (can’t search for partial metric names)



# Graph Generator Tip

- Undoing Functions
  - Select a metric in the Graph Data
  - Click Undo Function one or more times to remove outer-most function(s) (like peeling an onion)



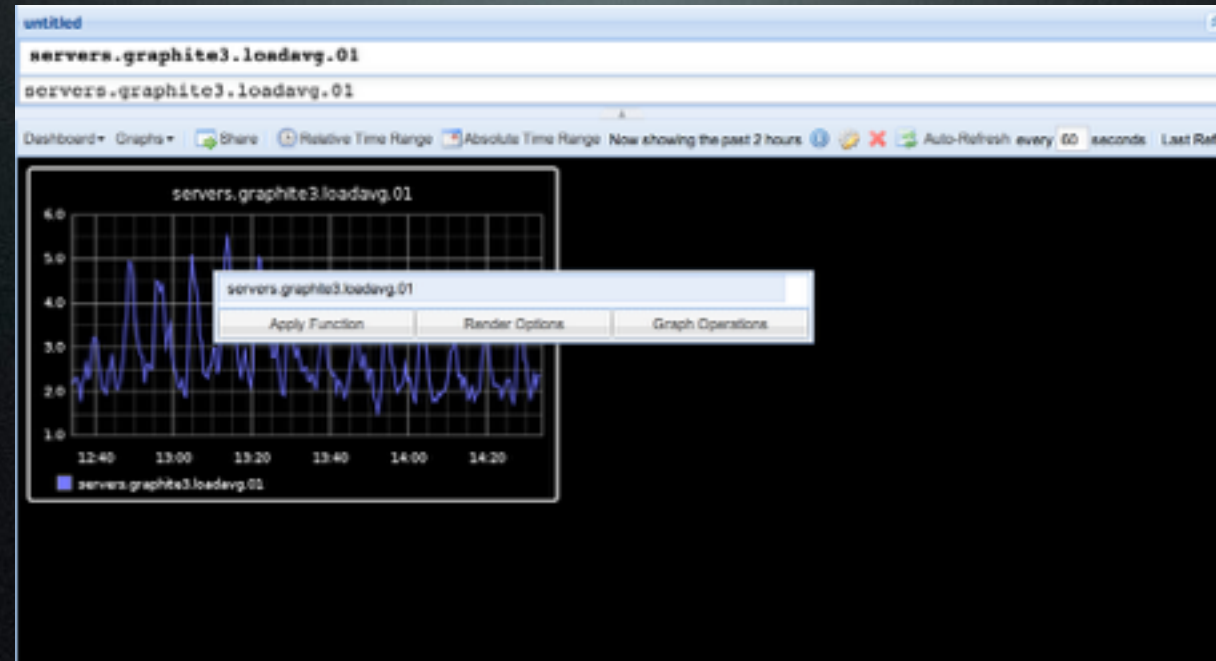
# Dashboard Generator

- Pros:
  - Generate dashboards of multiple graphs
  - Save dashboards for later review
  - Apply/remove functions on metrics
- Cons:
  - Can't easily remove a metric from a graph
  - No per-user/group permissions for dashboards
  - Dashboards saved in SQLite\*

\* There is support for storing dashboards in PostgreSQL, but it currently isn't supported at AWeber.  
– Important when setting up clusters of Graphite servers so that dashboards are shared across servers




# Dashboard Generator



# Dashboard Generator Tip

- Change UI to Tree view for easier navigation



Need to remember to refresh the page after changing the view (did you remember to save everything you were working on??)

# Dashboard Generator Tip

- Configuration & auto-refresh buttons




- Keyboard shortcuts

Keyboard Shortcuts	
Shortcut	Action
Ctrl-z	Toggle visibility of the graph area toolbar
Ctrl-space	Toggle visibility of the navigation panel
Alt-x	Remove all graphs from the graph area
Alt-s	Save the current dashboard
<i>The following shortcuts are for the Completer UI mode only</i>	
Alt-Enter	Add all matching metrics to the graph area
Alt-Backspace	Remove all matching metrics from the graph area
Shift-space	Put keyboard focus in the completer field

# Dashboard Generator Tip

- Hide the tree view



- Open the tree view



Cool! Now What?



# Quick Gallery

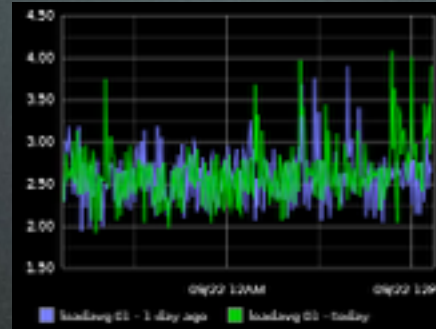
# Pie Charts!?



[http://graphite.colo.lair/render/?target=aliasByNode\(servers.graphite3.loadavg.{01,05,15},2,3\)&graphType=pie](http://graphite.colo.lair/render/?target=aliasByNode(servers.graphite3.loadavg.{01,05,15},2,3)&graphType=pie)

- \* Uses the 'aliasByNode' function to define the legend aliases
  - Note that both the 2nd & 3rd position of the metric name are used to create the alias
- \* 'graphType' is used to make it a pie chart
  - 'line' & 'pie' are the only options available currently
  - A 'bar' option has been added as a patch, but not sure it's merged into production yet upstream

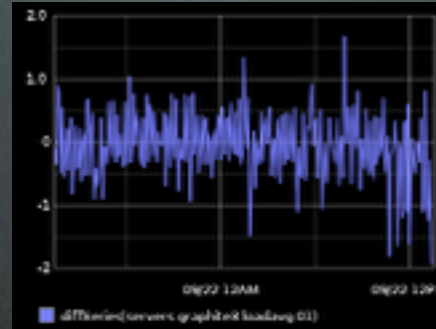
# Time Shifted Comparisons



[http://graphite.colo.lair/render/?target=alias\(timeShift\(servers.graphite3.loadavg.01%2C%221d%22\),%22loadavg%2001%20-%201%20day%20ago%22\)&target=alias\(servers.graphite3.loadavg.01,%22loadavg%2001%20-%20today%22\)](http://graphite.colo.lair/render/?target=alias(timeShift(servers.graphite3.loadavg.01%2C%221d%22),%22loadavg%2001%20-%201%20day%20ago%22)&target=alias(servers.graphite3.loadavg.01,%22loadavg%2001%20-%20today%22))

- \* Uses the 'timeShift' function to overlay a metric from a day ago with the same metric for today
- \* Great for comparing historic data from a comparative time period with current data

# Difference Between Series




[http://graphite3 colo.lair/render/?target=diffSeries\(timeShift\(servers.graphite3.loadavg.01,%221d%22\),servers.graphite3.loadavg.01\)](http://graphite3 colo.lair/render/?target=diffSeries(timeShift(servers.graphite3.loadavg.01,%221d%22),servers.graphite3.loadavg.01))

Shows the difference in 1-minute load on graphite3 today vs. the same times yesterday.

- \* Potentially useful in defining checks that are more adaptive
- \* Allows for monitoring amount of change from historic values rather than setting specific thresholds

# Most Deviant



[http://graphite3.colo.lair/render/?target=mostDeviant\(5,servers.\\*.loadavg.01\)](http://graphite3.colo.lair/render/?target=mostDeviant(5,servers.*.loadavg.01))

Shows the 5 most deviant (outlier) metrics from a grouping.

In this case it's the top 5 1-minute load averages across all servers.



# Resources

- Jason Dixon's blog (tips, tricks, tools)
  - <http://obfuscureity.com/Tags/Graphite>
- Graphite API Reference
  - <https://graphite.readthedocs.org/en/latest/functions.html>
- Clustering Graphite
  - <http://bitprophet.org/blog/2013/03/07/graphite/>

Questions?